

C1 - Assignment 1 Report: Sparse Matrices.

Student Number: 1894945

November 5, 2018

C1 - Assignment 1 Report

Student Number: 1894945

Contents

1	Introduction	3
1.1	Well-posed, direct problems	3
1.2	Numerical Methods	3
1.3	Linear systems	4
1.3.1	Splitting Methods	4
1.3.2	Gauss-Seidel Method	4
2	Implementation	5
2.1	Construction of the Gauss-Seidel function	5
2.2	Testing correctness	6
2.3	Running the program	6
2.4	Results of the tests	6
2.4.1	Varying the matrix size	6
2.4.2	Varying δ	7
2.4.3	varying λ	7
3	Conclusive remarks	7
3.1	Memory	7
3.2	Performances	7
3.2.1	Possible solutions	7

Abstract

1 Introduction

1.1 Well-posed, direct problems

The problems that will be addressed in the following are always representable in the form:

$$F(x, d) = 0 \tag{1}$$

where x represents the unknown, d the set of data from which the solution depends on and F the functional relation between x and d . Such types of problem are called *direct problems* ([1]).

Definition 1. Let D be the set of admissible data, i.e. the set of data for which problem (1) admits a unique solution x . Let $d \in D$ and denote by δd a perturbation such that $d + \delta d \in D$ and by δx the corresponding change in the solution, in such a way that

$$F(d + \delta d, x + \delta x) = 0$$

Then the solution x depends continuously on the data d if

$$\exists \eta_0(d), \exists K_0(d)$$

such that:

$$\|\delta d\| \leq \eta_0(d) \implies \|\delta x\| \leq K_0(d) \|\delta d\|$$

If d is admissible for (1) and if the same problem admits a unique solution x continuously depending on the data d , then the problem is said to be *well-posed* or *stable*. Whenever the aforementioned properties are not satisfied, the problem is said to be *ill-posed*.

1.2 Numerical Methods

In the following, it will always be assumed that problem (1) is well-posed. A numerical method for the approximate solution of the aforementioned equation consists in a sequence of approximate problems:

$$F_n(x_n, d_n) = 0 \quad n \geq 1 \tag{2}$$

with the underlying expectation that $x_n \rightarrow x$ as $n \rightarrow \infty$, i.e. the approximate solution converges to the exact one.

Definition 2. Consider the problem

$$F_n(x_n, d_n) = 0, \quad n \in \mathbb{N}$$

and denote D_n the set of admissible data for this problem. Then, the numerical method F_n is stable if its solution x_n depends continuously on the data d_n , for all admissible data $d_n \in D_n$.

Definition 3. The numerical method (2) is convergent iff

$$\forall \epsilon > 0, \exists n_\epsilon, \exists \delta(n_\epsilon) \mid \forall n > n_\epsilon, \forall \delta d_n : \|x(d) - x_n(d + \delta d_n)\| < \epsilon$$

where δd_n a perturbation of d_n , d_n is an admissible datum for the n^{th} approximate problem, $x_n(d + \delta d_n)$ the corresponding solution of it and $x(d)$ the solution of the exact problem.

1.3 Linear systems

Consider the following linear system:

$$Ax = b$$

where $A \in \mathbb{R}^{n \times n}$, $x \in \mathbb{R}^n$ and $b \in \mathbb{R}^n$. It is evident that such a problem can be represented in the form (1) as follows:

$$F(x, (A, b)) = 0$$

Assuming A to be non singular, in order to obtain the solution x of the given problem, one should generally invert A , so that:

$$x = A^{-1}b$$

The standard *direct* method for inverting a generic dense matrix is given by computing its LU decomposition [2]. The general computational cost for such a procedure is $\mathcal{O}(n^3)$: the method becomes impractical if A is large or sparse, as it happens in the present assignment.

A naive implementation of a direct method to invert a sparse matrix is a computational waste, both in memory and compute time. One would rather not store the zeros in a sparse matrix, nor multiply by them. Certainly, a direct method can be adapted to do neither of these things, leading to the *sparse direct methods*. Alternatively, one can rely on the so called *iterative methods*. Iterative methods are a class of matrix inversion techniques depending only on the calculation of matrix-vector multiplications, which are typically relatively easy to implement in the case of sparse matrices [2].

The performance of an iterative method is generally evaluated by the number of iterations required to converge to a sufficiently accurate approximation of the actual solution. A generic sparse-matrix-vector multiplication costs $\mathcal{O}(kn)$, k being the average number of non-zero elements in each. Hence, ideally, one desires an iterative method whose number of iterations required to converge is either independent of n or scales sublinearly with respect to n , in order to be competitive with direct methods, without relying on the sparsity of the matrix.

1.3.1 Splitting Methods

In order to implement an iterative method, it is necessary to generate a sequence of approximations $\{x^{(k)}\}$ to the solution. In the case of a splitting method for the problem $Ax = b$, the sequence is obtained as follows:

$$Px^{(k+1)} = Nx^{(k)} + b$$

or, equivalently:

$$x^{(k+1)} = x^{(k)} + P^{-1}r^{(k)}$$

where $A = P - N$ is the *splitting matrix*, P is the preconditioner matrix and $r^{(k)} = b - Ax^{(k)}$. In particular, the inversion of P , which has to be non singular, has not cost more than $\mathcal{O}(n^2)$ operations, in order to have $\mathcal{O}(n^2)$ operations for the whole method (the left multiply with N is always this much).

1.3.2 Gauss-Seidel Method

Let the decomposition of A be redefined as $A = D - (E + F)$, where D is the diagonal of A , $-E$ and $-F$ are the upper-triangular and lower-triangular components of $A - D$.

The Gauss-Seidel method is defined by setting $P_G = D - E$ and $N_G = F$. Clearly, inverting P_G costs $\mathcal{O}(n^2)$. It can be proven (see [2], [3]) that, if $A \in \mathbb{R}^{n \times n}$ is symmetric and positive definite, then the Gauss-Seidel

iteration converges for any $x^{(0)}$.¹

Taking advantage of the triangular form of P , the updating of the i^{th} component of x at the $(k+1)^{th}$ iteration can be written in the following form:

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left[b_i - \sum_{j=1}^{i-1} a_{ij} x_j^{(k+1)} - \sum_{j=i+1}^n a_{ij} x_j^{(k)} \right] \quad i = 1, \dots, n \quad (3)$$

Hence, in the Gauss-Seidel method, at the $(k+1)^{th}$ step, the available values of $x^{(k+1)}$ are used to update the solution.

The algorithm is clearly terminated on the residual $r^{(k)}$, meaning that the iterations stop when

$$\|r^{(k)}\| < TOL$$

TOL being some error tolerance. It is, however, not clear with respect to which norm the previous condition has to be satisfied: even though the choice a norm is completely subjective, L_∞ and L_2 norms are the ones most commonly used (see [2]).

2 Implementation

The program has been built with the aim to minimise the memory used to store the matrices. Hence, following the instructions given in the assignment, the STL container *vector* has been used to store both the matrix entries and the corresponding column indexes. However, instead of using simple vectors, vectors of pointers have been used. Such a solution provides a more efficient storing of both entries and indexes, as it can be directly seen by comparing the memory allocated by the program with the case in which simple vectors are used.

2.1 Construction of the Gauss-Seidel function

The core of the whole program is clearly represented by the function performing the Gauss-Seidel algorithm. Such a function has been chosen to be a member function and has been built to accept 7 parameters: the initial guess x_0 , the vector to invert against b , the tolerance required by the method (*tol*), the number of iterations after which a check for stagnation has to be made (*itCheck*), two files on which the residual error at every iteration and the solution are printed respectively, the maximum number of iteration the user is willing to wait for the algorithm to converge (*MaxIter*).

Looking at 3, it is clear that the elements of x_0 can be overwritten as they are computed in the algorithm and only one storage vector is needed: x_0 has then been passed by reference to the function, so that no copy of it has to be produced and efficiency is improved.

In case the matrix is not a square matrix and the sizes of x_0 and b do not equal the row size column and row numbers of the matrix, the program exits with an appropriate warning to the user. If this is not the case, the iterations start. The program runs until either the residual error becomes smaller than the given tolerance or the maximum number of iterations *MaxIter* has been reached. In order to make sure that the program has not stagnated, every *itCheck* iterations, the current residual error, in the L_∞ norm, is compared with the one computed and stored *itCheck* iterations before: if the residual has not decreased, the program exits with an appropriate stagnation warning to the user.

The sudo code for the implementation of the algorithm can be found either in Ref.[2] or on Wikipedia and will not be reported here.

¹Clearly, in order to improve convergence of the algorithm, an optimal choice of the *initial guess* $x^{(0)}$, has to be made.

2.2 Testing correctness

In order to check the correctness of the implementation of the Gauss-Seidel algorithm, the tests requested by the assignment have been performed: for this reason, a non-member function `Gauss_Seidel_test` has been built. Such a function takes 6 arguments: the size of the matrix A that has to be initialised to perform the test, the parameters δ and λ necessary to the implementation of the same tests, the fixed tolerance tol , the number of iterations after which a check for stagnation has to be made (*itCheck*), the maximum number of iteration the user is willing to wait for the algorithm to converge (*MaxIter*). The last 3 parameters will be used by the member function `Gauss_Seidel` previously introduced, within the function `Gauss_Seidel_test`.

The vectors x_0 , b , ω and D and the parameter $a = 4(\delta - 1)$ necessary to construct the matrix A are initialised by this function. The entries of A are added by means of the member function *addEntry*.

The files on which the residual error at every iteration and the corresponding final solution are printed on, are identified by the value of δ , λ by the strings *Residual* or *Solution* respectively and by the size of the matrix the test has been performed on.

The function prints on the terminal the error between the exact and the found solution, indexing it through the value of λ , δ and the size N of matrix A .

2.3 Running the program

The program is run through the Makefile provided. The current optimisation is set to `-Ofast` and all following tests have been performed with this optimisation choice. Every possible error or warning has been explicitly checked by means of the instruction `-Wall -Wfatal-errors -pedantic` (that can still be found commented out in the Makefile) before proceeding to performing the different tests: no warnings appeared.

The program generates a fixed number of output files, in which the results of the tests have been stored. Some of these files are used to produce plots by means of two different plotscripts.

2.4 Results of the tests

In the following, the results of the more significant tests performed will be reported and briefly discussed.

2.4.1 Varying the matrix size

The `Gauss_Seidel_test` function has been show to work for the case of $N = 100$, $N = 1000$ and $N = 10000$, as required by the assignment. The time T_{10000} required to reach convergence in the case $N = 10000$ has been explicitly computed trough the following instructions:

```
auto start = std::chrono::high_resolution_clock::now();

auto finish = std::chrono::high_resolution_clock::now();
std::chrono::duration<double> elapsed = finish - start;
std::cout<< "Elapsed time: " << elapsed.count() << "s\n";
```

and found to be $T_{10000} \approx 83$ seconds. The method has thus been shown to work in a reasonably short time, so that the program can be regarded as being fairly efficient.

2.4.2 Varying δ

2.4.3 varying λ

3 Conclusive remarks

3.1 Memory

Even though the memory for storing the matrix entries and indeces is dynamically allocated, no explicit deletion of it has been performed: this is due to the
From now on everything with -g and -O1 optimisation.
Controlling memory leaks.

```
$ valgrind --leak-check=yes ./sparsematrix
```

```
==4200== Memcheck, a memory error detector
==4200== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==4200== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright info
==4200== Command: ./sparsematrix
==4200==
==4200==
==4200== HEAP SUMMARY:
==4200==    in use at exit: 0 bytes in 0 blocks
==4200==   total heap usage: 1,252,428 allocs, 1,252,428 frees, 989,616,540 bytes allocated
==4200==
==4200== All heap blocks were freed -- no leaks are possible
==4200==
==4200== For counts of detected and suppressed errors, rerun with: -v
==4200== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

No memory lost. Everything deleted correctly.

3.2 Performances

Also did:

```
$ valgrind --tool=callgrind ./sparsematrix
$ kcache-grind
```

First command analyses the performances in the terms of load distribution. kcache-grind to visualise the visualise the load distribution results.
INCOLLA FOTO

3.2.1 Possible solutions

References

- [1] A. Quateroni, R. Sacco, F. Saleri; *Numerical Mathematics*, Vol.37, Springer Verlag, (2007).
- [2] T. Grafke; *Scientific Computing*, Lecture Notes, University of Warwick, (2018).
- [3] W. Hackbush; *Iterative Solution of Large Sparse Systems of Equations*, Springer-Verlag, New York, (1994).